

1. Use all-subsets regression.

(a) Report the variables in the best model of each size.

```
### The model.matrix lets us specify the regression formula we want
### to use, and outputs the corresponding model matrix
data.matrix = model.matrix(Ozone ~ .^2, data = data.train)

### We also need the response variable (i.e. alcohol)
Y.train = data.train$Ozone

### Now we can run all.subsets. There are a couple of extra inputs
### here. nvmax is the largest number of variables we are willing
### to include. 30 seems like plenty. intercept specifies whether we
### want regsubsets() to add an intercept term. Since model.matrix()
### already adds an intercept, we don't want regsubsets() to add
### another one.
all.subsets = regsubsets(x = data.matrix, y = Y.train, nvmax = 30)

### The output of regsubsets isn't really useful. We need to run the
### summary() function on it to get useful information
info.subsets = summary(all.subsets)

### The output of summary contains an array with columns corresponding
### to predictors and rows corresponding to model sizes. This array
### tells us which variables are included at each size.
all.subsets.models = info.subsets$which
all.subsets.models = all.subsets.models[, -1]
all.subsets.models
dim(all.subsets.models)

### We can get the AIC and BIC of each of these models by re-fitting
### the models and running extractAIC(). The extractAIC() function
### has an input called k, which is the coefficient on the penalty term.
```

```
> all.subsets.models
(Intercept) Solar.R Wind Temp Twcp Twrat Solar.R:Wind Solar.R:Temp Solar.R:Twcp Solar.R:Twrat Wind:Temp Wind:Twcp Wind:Twrat Temp:Twcp
1 FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE TRUE FALSE FALSE FALSE FALSE
2 FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE TRUE FALSE FALSE FALSE FALSE
3 FALSE FALSE FALSE FALSE FALSE FALSE FALSE TRUE FALSE FALSE TRUE FALSE FALSE FALSE FALSE
4 FALSE FALSE FALSE FALSE FALSE FALSE TRUE FALSE FALSE FALSE TRUE FALSE FALSE FALSE FALSE
5 FALSE FALSE FALSE TRUE FALSE TRUE FALSE FALSE FALSE TRUE FALSE FALSE FALSE FALSE FALSE
6 FALSE FALSE TRUE TRUE FALSE TRUE FALSE FALSE FALSE TRUE FALSE FALSE FALSE FALSE FALSE
7 FALSE FALSE FALSE TRUE FALSE TRUE TRUE TRUE TRUE TRUE FALSE FALSE FALSE FALSE FALSE
8 FALSE TRUE FALSE FALSE FALSE TRUE TRUE TRUE FALSE TRUE TRUE FALSE FALSE TRUE FALSE
9 FALSE FALSE FALSE FALSE FALSE TRUE TRUE TRUE TRUE TRUE TRUE TRUE FALSE TRUE TRUE
10 FALSE FALSE FALSE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE FALSE TRUE TRUE
11 FALSE TRUE FALSE TRUE TRUE TRUE FALSE TRUE TRUE TRUE TRUE FALSE TRUE TRUE TRUE
12 FALSE TRUE TRUE TRUE TRUE TRUE FALSE TRUE TRUE TRUE TRUE FALSE TRUE TRUE TRUE
13 FALSE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE FALSE TRUE FALSE TRUE

Temp:Twrat Twcp:Twrat
1 FALSE FALSE
2 FALSE TRUE
3 FALSE TRUE
4 TRUE TRUE
5 TRUE TRUE
6 TRUE TRUE
7 TRUE TRUE
8 TRUE TRUE
9 TRUE TRUE
10 TRUE TRUE
11 TRUE TRUE
12 TRUE TRUE
13 TRUE TRUE
```

(b) Compute BIC on each of these models and **report the BIC values for the models.**

```
n.models = nrow(all.subsets.models) # Number of candidate models
all.AICs = rep(0, times = n.models) # Container to store AICs
all.BICs = all.AICs # Copy all.AICs to get a container for BICs

for(i in 1:n.models){
  ### We can actually supply a model matrix and response vector
  ### to lm, without using a data frame. Remember that our model matrix
  ### already has an intercept, so we need to make sure lm doesn't
  ### include another one. We do this by including -1 in the right side
  ### of the model formula.
  this.data.matrix = data.matrix[,all.subsets.models[i,]]
  fit = lm(Y.train ~ this.data.matrix - 1)

  ### Get the AIC using extractAIC(). This function takes a regression
  ### model as input, as well as (optionally) an input called k, which
  ### specifies the penalty on the number of variables in our model.
  ### The AIC value is in the second component of the output object.
  this.AIC = extractAIC(fit)[2]
  all.AICs[i] = this.AIC

  ### Get the BIC using extractAIC(). This time, we need to set k equal
  ### to the log of the number of observations used to fit our model
  this.BIC = extractAIC(fit, k = log(n.train))[2]
  all.BICs[i] = this.BIC
}

> all.BICs
[1] 527.0872 514.9455 513.8233 522.6431 501.1093 497.0320 501.0831 500.8298 493.0466 496.3106 498.2233 489.8231 494.1423
```

(c) Identify the best model. **What variables are in it?**

```
> which.min(all.BICs)
[1] 12
```

12th model includes "Solar.R", "Wind", "Temp", "TWcp", "TWrat", "Solar.R:Temp", "Solar.R:TWcp", "Solar.R:TWrat", "Wind:TWcp", "Temp:TWcp", "Temp:TWrat", and "TWcp:TWrat"

2. Use the hybrid stepwise algorithm that is the default in the step() function. **Report the model that it chooses as “best.”**

```
#(b) Compute BIC on each of these models and report the BIC values for the
#models.
```

```
#####Use stepwise
```

```
fit.start = lm(Ozone ~ 1, data = data.train)
```

```
fit.end = lm(Ozone ~ .^2, data = data.train)
```

```
step.AIC = step(fit.start, list(upper = fit.end), k = 2)
```

```
step.BIC = step(fit.start, list(upper = fit.end), k = log(n.train), trace = 0)
```

```
pred.step.AIC = predict(step.AIC, data.valid)
```

```
pred.step.BIC = predict(step.BIC, data.valid)
```

```
err.step.AIC = get.MSPE(Y.valid, pred.step.AIC)
```

```
err.step.BIC = get.MSPE(Y.valid, pred.step.BIC)
```

```
> summary(step.BIC)
```

Call:

```
lm(formula = Ozone ~ TWrat + Temp + Solar.R + TWrat:Solar.R +
    TWrat:Temp, data = data.train)
```

Residuals:

```
      Min       1Q   Median       3Q      Max
-37.563 -12.930  -3.048  10.033  46.888
```

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	-1.584e+02	3.400e+01	-4.658	1.31e-05	***
TWrat	1.068e+01	4.508e+00	2.370	0.020312	*
Temp	2.555e+00	4.351e-01	5.872	1.03e-07	***
Solar.R	-1.705e-01	6.466e-02	-2.636	0.010129	*
TWrat:Solar.R	3.060e-02	7.479e-03	4.091	0.000105	***
TWrat:Temp	-1.613e-01	4.965e-02	-3.248	0.001724	**

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 18.88 on 77 degrees of freedom

Multiple R-squared: 0.7313, Adjusted R-squared: 0.7138

F-statistic: 41.91 on 5 and 77 DF, p-value: < 2.2e-16

The model that use TWrat, Temp, Solar.R, TWrat:Solar.R, and TWrat:Temp as variables

3. Use 10-fold CV to estimate the MSPE for the stepwise model selection process. That is,

##(c) Identify the best model. What variables are in it?

```
set.seed(2928893)|
### First we need to set the number of folds
K = 10

### Construct folds
### Don't attach fold labels to dataset because we would just have
### to remove this later
n = nrow(data)
n.fold = n/K # Approximate number of observations per fold
n.fold = ceiling(n.fold)
ordered.ids = rep(1:10, each = n.fold)
ordered.ids = ordered.ids[1:n]
fold.ids = shuffle(ordered.ids)

### Create a container to store CV MSPEs
### One column per model, and one row per fold
CV.models = c("stepwise.AIC", "stepwise.BIC")
errs.CV = array(0, dim = c(K,length(CV.models)))
colnames(errs.CV) = CV.models

for(i in 1:K){
  print(paste0(i, " of ", K))

  ### Construct training and validation sets by either removing
  ### or extracting the current fold.
  ### Also, get the response vectors
  data.train = data[fold.ids != i,]
  data.valid = data[fold.ids == i,]
  Y.train = data.train$Ozone
  Y.valid = data.valid$Ozone

  #####
  ### Stepwise selection via AIC and BIC ###
  #####

  fit.start = lm(Ozone ~ 1, data = data.train)
  fit.end = lm(Ozone ~ .^2, data = data.train)

  ### These functions will run several times each. We don't need
  ### to print out all the details, so set trace = 0.
  step.AIC = step(fit.start, list(upper = fit.end), k=2,
                  trace = 0)
  step.BIC = step(fit.start, list(upper = fit.end), k = log(n.train),
                  trace = 0)
  print(summary(step.BIC))

  pred.step.AIC = predict(step.AIC, data.valid)
  pred.step.BIC = predict(step.BIC, data.valid)
```



```

err.step.AIC = get.MSPE(Y.valid, pred.step.AIC)
err.step.BIC = get.MSPE(Y.valid, pred.step.BIC)

### Store errors in errs.CV, which has two dimensions, so
### we need two indices
errs.CV[i, "stepwise.AIC"] = err.step.AIC
errs.CV[i, "stepwise.BIC"] = err.step.BIC
}

```

```

> errs.CV
      stepwise.AIC stepwise.BIC
[1,]      1445.724      1379.311
[2,]      2929.978      3469.082
[3,]      1787.337      1720.089
[4,]      1320.081      1320.081
[5,]      2407.516      2407.516
[6,]      2302.269      2263.954
[7,]      2003.969      2066.328
[8,]      1529.626      1544.587
[9,]      2007.500      2007.500
[10,]     2224.739      2563.569

```

4th model shows the smallest BIC with 1320.081.