

5 Exercises

Application

Refer to the Air Quality data described previously, and the analyses we have done with Ozone as the response variable, and the five explanatory variables (including the two engineered features).

Use Neural Nets (NN) to model the relationship between Ozone and all five explanatories as specified below.:

1. Create a matrix of the five explanatory variables. Rescale each of them to lie between 0 and 1 and save them as another matrix. **Print a summary() of each object to confirm that you have scaled properly.**

```
1 # Title: STAT 452 Exercise 8 L12Q12
2 # Author: Injun Son
3 # Date: October 31, 2020
4
5 library(dplyr)
6 library(MASS) # For ridge regression
7 library(glmnet) # For LASSO
8 source("Helper Functions.R")
9 data = na.omit(airquality[, 1:4])
10 data$TWcp = data$Temp*data$Wind
11 data$TWrat = data$Temp/data$Wind
12
13
14 # 1. Create a matrix of the five explanatory variables. Rescale each of them to lie between
15 # 0 and 1 and save them as another matrix. Print a summary() of each object to
16 # confirm that you have scaled properly.
17
18 ### 75%/25% split into training and validation sets
19 n = nrow(data)
20 n.train = floor(n*0.75)
21 n.valid = n - n.train
22
23 inds = c(rep(1, times = n.train), rep(2, times = n.valid))
24 inds.rand = inds[sample.int(n)]
25
26 data.train = data[inds.rand == 1,]
27 X.train.raw = data.train[,-5]
28 Y.train = data.train[,5]
29
30 data.valid = data[inds.rand == 2,]
31 X.valid.raw = data.valid[,-5]
32 Y.valid = data.valid[,5]
33
34
35
36
37
38
39
40 rescale <- function(x1,x2){
41   for(col in 1:ncol(x1)){
42     a <- min(x2[,col])
43     b <- max(x2[,col])
44     x1[,col] <- (x1[,col]-a)/(b-a)
45   }
46   x1
47 }
48
49 X.train = rescale(X.train.raw, X.train.raw)
50 X.valid = rescale(X.valid.raw, X.train.raw) # Be careful with the order
```

```
> summary(X.train)
      Ozone      Solar.R      Wind      Temp      Twrat
Min.   :0.0000   Min.   :0.0000   Min.   :0.0000   Min.   :0.0000   Min.   :0.00000
1st Qu.:0.1018   1st Qu.:0.3104   1st Qu.:0.2772   1st Qu.:0.2763   1st Qu.:0.07554
Median :0.1796   Median :0.6269   Median :0.4022   Median :0.5000   Median :0.11994
Mean   :0.2402   Mean   :0.5430   Mean   :0.4151   Mean   :0.4807   Mean   :0.16533
3rd Qu.:0.2904   3rd Qu.:0.7661   3rd Qu.:0.5000   3rd Qu.:0.6579   3rd Qu.:0.20358
Max.   :1.0000   Max.   :1.0000   Max.   :1.0000   Max.   :1.0000   Max.   :1.00000

> summary(X.valid)
      Ozone      Solar.R      Wind      Temp      Twrat
Min.   :0.02994   Min.   :0.05199   Min.   :0.02717   Min.   : -0.05263   Min.   :0.001665
1st Qu.:0.08683   1st Qu.:0.32798   1st Qu.:0.27038   1st Qu.: 0.40789   1st Qu.:0.084249
Median :0.17964   Median :0.60245   Median :0.43478   Median : 0.57895   Median :0.117182
Mean   :0.26369   Mean   :0.54587   Mean   :0.41557   Mean   : 0.53571   Mean   :0.178850
3rd Qu.:0.42216   3rd Qu.:0.75306   3rd Qu.:0.54348   3rd Qu.: 0.71711   3rd Qu.:0.220365
Max.   :0.68263   Max.   :0.89297   Max.   :0.87500   Max.   : 0.92105   Max.   :0.797665
```

2. Fit 4 NNs nets using only Temp and Wind, using each combination of 2 and 6 hidden nodes with 0.001 and 1 shrinkage; i.e., (2,0.001), (2,1), (6,0.001), (6,1)

(a) Refit each one manually 20 times or more and compute the sMSE each time.

i. **Report the sMSE for the optimal fit for each model.**

(2, 0.001)

```
70 ##### (2, 0.001)
71
72 n.hidden = 2
73 shrink = 0.001
74
75 fit.nnet = nnet(y = Y.train, x = X.train[,c(3,4)], linout = TRUE, size = n.hidden,
76               decay = shrink, maxit = 500)
77
78 n.nnets = 20 # Number of times to re-fit
79 ### Container for SSEs
80 all.SSEs = rep(0, times = 20)
81 all.MSEs = rep(0, times = 20)
82 all.nnets = list(1:20)
83 for(i in 1:n.nnets){
84   ### Fit model. We can set the input "trace" to FALSE to suppress the
85   ### printed output from the nnet() function.
86   this.nnet = nnet(y = Y.train, x = X.train[,c(3,4)], linout = TRUE, size = n.hidden,
87                 decay = shrink, maxit = 500, trace = FALSE)
88
89   ### Get the model's SSE
90   this.SSE = this.nnet$value
91   this.MSE = this.nnet$value / nrow(X.train)
92
93   ### Store results. We have to use double square brackets when storing or
94   ### retrieving from a list.
95   all.SSEs[i] = this.SSE
96   all.MSEs[i] = this.MSE
97   all.nnets[[i]] = this.nnet
98 }
99 all.MSEs
100
101 ind.best = which.min(all.SSEs)
102 fit.nnet.best = all.nnets[[ind.best]]
103
104 x1 <- seq(from=2.3, to=20.7, by=.5)
105 x2 = seq(from=57, to=97, by=.5)
106 xy1 <- data.frame(expand.grid(wind=x1, temp=x2))
107
108 pred2 <- predict(fit.nnet, newdata=rescale(xy1, X.valid.raw[,c(3,4)]))
109 surface2 = matrix(pred2, nrow=length(x1))
```

```

110
111
112 open3d()
113 persp3d(x = x1, y = x2,
114         z = surface2, col = "orange", xlab="Wind", ylab="Temp",
115         zlab="Ozone")
116 points3d(data$Ozone ~ data$Wind + data$Temp, col="blue")
117
118 > all.MSEs
[1] 56.16532 56.33156 93.24751 109.80589 58121.62992 43.01570 109.80538 40.53980 58120.69365
[10] 105.25589 25434.99372 109.80613 58125.31099 58120.17503 31.22951 43.01570 99.33437 56.16532
[19] 109.80778 79.74207
119 ##### (2, 1)
120 n.hidden = 2
121 shrink = 1
122
123 fit.nnet = nnet(y = Y.train, x = X.train[,c(3,4)], linout = TRUE, size = n.hidden,
124               decay = shrink, maxit = 500)
125
126 n.nnets = 20 # Number of times to re-fit
127 ### Container for SSEs
128 all.SSEs = rep(0, times = 20)
129 all.MSEs = rep(0, times = 20)
130 all.nnets = list(1:20)
131 for(i in 1:n.nnets){
132   ### Fit model. We can set the input "trace" to FALSE to suppress the
133   ### printed output from the nnet() function.
134   this.nnet = nnet(y = Y.train, x = X.train[,c(3,4)], linout = TRUE, size = n.hidden,
135                  decay = shrink, maxit = 500, trace = FALSE)
136
137   ### Get the model's SSE
138   this.SSE = this.nnet$value
139   this.MSE = this.nnet$value / nrow(X.train)
140
141   ### Store results. We have to use double square brackets when storing or
142   ### retrieving from a list.
143   all.SSEs[i] = this.SSE
144   all.MSEs[i] = this.MSE
145   all.nnets[[i]] = this.nnet
146 }
147 all.MSEs
148 ind.best = which.min(all.SSEs)
149 fit.nnet.best = all.nnets[[ind.best]]
150
151 x1 <- seq(from=2.3, to=20.7, by=.5)
152 x2 = seq(from=57, to=97, by=.5)
153 xy1 <- data.frame(expand.grid(wind=x1, temp=x2))
154 |
155 pred2 <- predict(fit.nnet, newdata=rescale(xy1, x.valid.raw[,c(3,4)]))
156 surface2 = matrix(pred2, nrow=length(x1))
157
158
159 open3d()
160 persp3d(x = x1, y = x2,

```

```

160 persp3d(x = x1, y = x2,
161         z = surface2, col = "orange", xlab="Wind", ylab="Temp",
162         zlab="Ozone")
163 points3d(data$Ozone ~ data$Wind + data$Temp, col="blue")
164
165 > all.MSEs
[1] 1721.232 1621.757 1137.290 1721.232 1954.603 1137.290 1137.290 1721.232 1137.290 1721.232 1621.757 1621.757 1137.290
[14] 1137.290 1721.232 1137.290 1721.232 1137.290 1137.290 1137.290

```

```

67 n.hidden = 6
68 shrink = 0.001
69
70 fit.nnet = nnet(y = Y.train, x = X.train, linout = TRUE, size = n.hidden,
71               decay = shrink, maxit = 500)
72
73 n.nnets = 20 # Number of times to re-fit
74 ### Container for SSEs
75 all.SSEs = rep(0, times = 20)
76 all.MSEs = rep(0, times = 20)
77 all.nnets = list(1:20)
78 for(i in 1:n.nnets){
79   ### Fit model. We can set the input "trace" to FALSE to suppress the
80   ### printed output from the nnet() function.
81   this.nnet = nnet(y = Y.train, x = X.train, linout = TRUE, size = n.hidden,
82                 decay = shrink, maxit = 500, trace = FALSE)
83
84   ### Get the model's SSE
85   this.SSE = this.nnet$value
86   this.MSE = this.nnet$value / nrow(X.train)
87
88   ### Store results. We have to use double square brackets when storing or
89   ### retrieving from a list.
90   all.SSEs[i] = this.SSE
91   all.MSEs[i] = this.MSE
92   all.nnets[[i]] = this.nnet
93 }
94 all.MSEs
95 ### Get the best model. We have to use double square brackets when storing or
96 ### retrieving from a list.
97 # ind.best = which.min(all.SSEs)
98 # fit.nnet.best = all.nnets[[ind.best]]
99 ind.best = which.min(all.MSEs)
100 fit.nnet.best = all.nnets[[ind.best]]
101
102 ### Get predictions and MSPE
103 pred.nnet = predict(fit.nnet.best, X.valid)
104 MSPE.nnet = get.MSPE(Y.valid, pred.nnet)

```

```

> all.MSEs
[1] 58119.579968 34.550258 34.433917 7.315829 15.705896 8.412257 58120.278681 8.862501 12.428587
[10] 8.813635 10.565405 58122.348502 58117.578898 58117.820407 13.542882 7.553607 6.763851 28.733807
[19] 12.141556 27.231437

```



```

222 ##### (6, 1)
223 n.hidden = 6
224 shrink = 1
225
226 fit.nnet = nnet(y = Y.train, x = X.train[,c(3,4)], linout = TRUE, size = n.hidden,
227               decay = shrink, maxit = 500)
228
229 n.nnets = 20 # Number of times to re-fit
230 ### Container for SSEs
231 all.SSEs = rep(0, times = 20)
232 all.MSEs = rep(0, times = 20)
233 all.nnets = list(1:20)
234 for(i in 1:n.nnets){
235   ### Fit model. We can set the input "trace" to FALSE to suppress the
236   ### printed output from the nnet() function.
237   this.nnet = nnet(y = Y.train, x = X.train[,c(3,4)], linout = TRUE, size = n.hidden,
238                 decay = shrink, maxit = 500, trace = FALSE)
239
240   ### Get the model's SSE
241   this.SSE = this.nnet$value
242   this.MSE = this.nnet$value / nrow(X.train)
243
244   ### Store results. We have to use double square brackets when storing or
245   ### retrieving from a list.
246   all.SSEs[i] = this.SSE
247   all.MSEs[i] = this.MSE
248   all.nnets[[i]] = this.nnet
249 }
250 all.MSEs
251 ### Get the best model. We have to use double square brackets when storing or
252 ### retrieving from a list.
253 # ind.best = which.min(all.SSEs)
254 # fit.nnet.best = all.nnets[[ind.best]]
255 ind.best = which.min(all.MSEs)
256 fit.nnet.best = all.nnets[[ind.best]]
257
258 x1 <- seq(from=2.3, to=20.7, by=.5)
259 x2 <- seq(from=57, to=97, by=.5)
260 xy1 <- data.frame(expand.grid(wind=x1, temp=x2))
261
262 #pred.nnet = predict(fit.nnet.best, X.valid)
263 # pred.nnet = predict(fit.nnet.best, X.valid)
264 # MSPE.nnet = get.MSPE(Y.valid, pred.nnet)
265 # pred2 <- predict(pred.nnet, newdata=rescale(xy1, X.valid.raw[,c(3,4)]))
266 pred2 <- predict(fit.nnet, newdata=rescale(xy1, X.valid.raw[,c(3,4)]))
267 surface2 = matrix(pred2, nrow=length(x1))
268
269
270 open3d()
271 persp3d(x = x1, y = x2,
272         z = surface2, col = "orange", xlab="Wind", ylab="Temp",
273         zlab="Ozone")
274 points3d(data$Ozone ~ data$Wind + data$Temp, col="blue")

```

```

> all.MSEs
[1] 3394.466 3889.075 3399.422 3394.466 3410.263 3402.865 3394.466 3401.277 3401.277 3398.238 3403.193 3403.194 3520.883
[14] 3520.882 3394.466 3394.471 3395.892 3400.495 3401.277 3395.892

```

ii. Comment on the stability of fits for different models. In other words, **which models were most/least consistent with the sMSE values produced by different fits?**

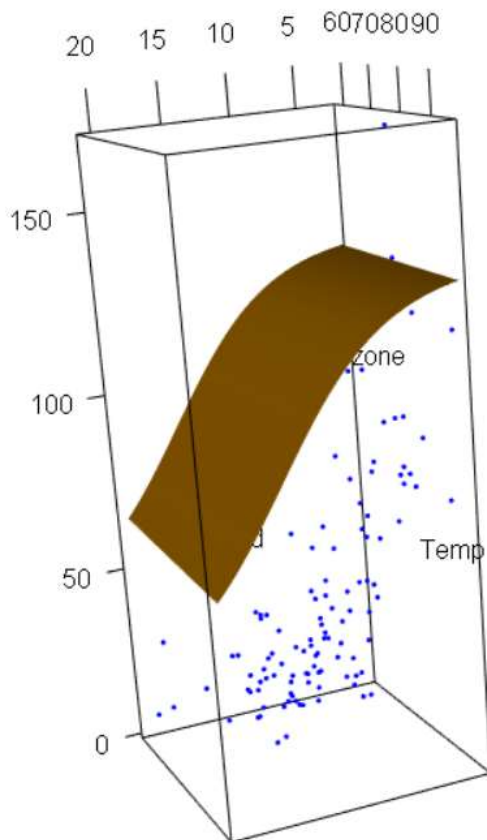
-> Model with 6 hidden nodes and 1 shrinkage has the most consistent sMSE.

-> Model with 2 hidden nodes and 0.001 shrinkage has the inconsistent sMSE.


(b) Make a 3-D plot of each model's fit.

i. Report a screenshot of each fit, rotated to roughly the same angle each time to show a good comparison of the fits. (I find it best to look down through the corner with low temp and high wind, so that the high ozone values are in the back.)


(2, 0.001)




(2, 1)



(6, 0.001)



(6, 1)



ii. **Comment separately on how increasing number of nodes or increasing shrinkage appears to affect the fits.**

➔ It looks like increasing shrinkage make more exact fits, but this looks like an over fit.