

2. Add three new models fitted to all 5 variables to the 10-fold CV comparison that has been used for LASSO, Ridge, and other methods: Full tree (CP=0), min-cv-pruned tree, and 1-se-pruned tree. Use the same folds as before.

```

120- get.folds = function(n, K) {
121-   ### Get the appropriate number of fold labels
122-   n.fold = ceiling(n / K) # Number of observations per fold (rounded up)
123-   fold.ids.raw = rep(1:K, times = n.fold) # Generate extra labels
124-   fold.ids = fold.ids.raw[1:n] # Keep only the correct number of labels
125-
126-   ### Shuffle the fold labels
127-   folds.rand = fold.ids[sample.int(n)]
128-   |
129-   return(folds.rand)
130- }
131-
132- # i=1
133- ### Number of folds
134- K = 10
135-
136- ### Construct folds
137- n = nrow(data) # Sample size
138- folds = get.folds(n, K)
139-
140- ### Create a container for MSPEs. Let's include ordinary least-squares
141- ### regression for reference
142- all.models = c("LS", "Hybrid", "Ridge", "LASSO-Min", "LASSO-lse", "GAM", "Full-tree", "Min-cv tree", "1-se tree")
143- all.MSPEs = array(0, dim = c(K, length(all.models)))
144- colnames(all.MSPEs) = all.models
145- ### Begin cross-validation
146- for(i in 1:K){
147-   ### Split data
148-   data.train = data[folds != i,]
149-   data.valid = data[folds == i,]
150-   n.train = nrow(data.train)
151-
152-   ### Get response vectors
153-   Y.train = data.train$Ozone
154-   Y.valid = data.valid$Ozone
155-
156-   # LS
157-   fit.ls = lm(Ozone ~ ., data = data.train)
158-   pred.ls = predict(fit.ls, newdata = data.valid)
159-   MSPE.ls = get.MSPE(Y.valid, pred.ls)
160-   all.MSPEs[i, "LS"] = MSPE.ls
161-
162-   #Hybrid Stepwise

```

```

163
164 fit.start = lm(Ozone ~ 1, data = data.train)
165 fit.end = lm(Ozone ~ ., data = data.train)
166
167 step.BIC = step(fit.start, list(upper = fit.end), k = log(n.train),
168                 trace = 0)
169
170 pred.step.BIC = predict(step.BIC, data.valid)
171
172 err.step.BIC = get.MSPE(Y.valid, pred.step.BIC)
173
174 all.MSPEs[i, "Hybrid"] = err.step.BIC
175
176
177
178 #ridge regression
179 lambda.vals = seq(from = 0, to = 100, by = 0.05)
180
181 fit.ridge = lm.ridge(Ozone ~ ., lambda = lambda.vals,
182                     data = data.train)
183
184 ind.min.GCV = which.min(fit.ridge$GCV)
185 lambda.min = lambda.vals[ind.min.GCV]
186
187 all.coefs.ridge = coef(fit.ridge)
188 coef.min = all.coefs.ridge[ind.min.GCV,]
189
190 matrix.valid.ridge = model.matrix(Ozone ~ ., data = data.valid)
191
192 ### Now we can multiply the data by our coefficient vector. The
193 ### syntax in R for matrix-vector multiplication is %*. Note that,
194 ### for this type of multiplication, order matters. That is,
195 ### A %*% B != B %*% A. Make sure you do data %*% coefficients.
196 ### For more information, see me in a Q&A session or, better still,
197 ### take a course on linear algebra (it's really neat stuff)
198 pred.ridge = matrix.valid.ridge %*% coef.min
199
200 ### Now we just need to calculate the MSPE and store it
201 MSPE.ridge = get.MSPE(Y.valid, pred.ridge)
202 all.MSPEs[i, "Ridge"] = MSPE.ridge
203
204 matrix.train.raw = model.matrix(Ozone ~ ., data = data.train)
205 matrix.train = matrix.train.raw[,-1]
206

```

```

207 ### LASSO
208 all.LASSOs = cv.glmnet(x = matrix.train, y = Y.train)
209
210 ### Get both 'best' lambda values using $lambda.min and $lambda.1se
211 lambda.min = all.LASSOs$lambda.min
212 lambda.1se = all.LASSOs$lambda.1se
213
214 ### Get the coefficients for our two 'best' LASSO models
215 coef.LASSO.min = predict(all.LASSOs, s = lambda.min, type = "coef")
216 coef.LASSO.1se = predict(all.LASSOs, s = lambda.1se, type = "coef")
217
218 ### Get which predictors are included in our models (i.e. which
219 ### predictors have non-zero coefficients)
220 included.LASSO.min = predict(all.LASSOs, s = lambda.min,
221                             type = "nonzero")
222 included.LASSO.1se = predict(all.LASSOs, s = lambda.1se,
223                             type = "nonzero")
224
225 matrix.valid.LASSO.raw = model.matrix(Ozone ~ ., data = data.valid)
226 matrix.valid.LASSO = matrix.valid.LASSO.raw[,-1]
227 pred.LASSO.min = predict(all.LASSOs, newx = matrix.valid.LASSO,
228                          s = lambda.min, type = "response")
229 pred.LASSO.1se = predict(all.LASSOs, newx = matrix.valid.LASSO,
230                          s = lambda.1se, type = "response")
231
232 ### Calculate MSPEs and store them
233 MSPE.LASSO.min = get.MSPE(Y.valid, pred.LASSO.min)
234 all.MSPEs[i, "LASSO-Min"] = MSPE.LASSO.min
235
236 MSPE.LASSO.1se = get.MSPE(Y.valid, pred.LASSO.1se)
237 all.MSPEs[i, "LASSO-1se"] = MSPE.LASSO.1se
238
239 ## GAM
240 fit.gam = gam(Ozone ~ s(Solar.R) + s(Wind) + s(Temp) + s(TWcp) + s(TWrat),
241               data = data.train)
242
243 pred.gam = predict(fit.gam, data.valid)
244 MSPE.gam = get.MSPE(Y.valid, pred.gam) # Our helper function
245 all.MSPEs[i, "GAM"] = MSPE.gam
246
247
248 # Full-tree
249 fit.tree = rpart(Ozone ~ ., data = data, cp=0)

```

```

250 fit.tree.pred = predict(fit.tree, data.valid)
251 MSPE.fit.tree = get.MSPE(Y.valid, fit.tree.pred)
252 all.MSPEs[i, "Full-tree"] = MSPE.fit.tree
253
254
255 # Min-cv tree
256 info.tree = fit.tree$cptable
257 info.tree
258 ind.min = which.min(info.tree[, "xerror"])
259 CP.min.raw = info.tree[ind.min, "CP"]
260 if(ind.min == 1){
261     ### If minimum CP is in row 1, store this value
262     CP.min = CP.min.raw
263 } else{
264     ### If minimum CP is not in row 1, average this with the value from the
265     ### row above it.
266
267     ### Value from row above
268     CP.above = info.tree[ind.min-1, "CP"]
269
270     ### (Geometric) average
271     CP.min = sqrt(CP.min.raw * CP.above)
272 }
273 fit.tree.min = prune(fit.tree, cp = CP.min)
274 fit.tree.min.pred = predict(fit.tree.min, data.valid)
275 MSPE.fit.tree.min = get.MSPE(Y.valid, fit.tree.min.pred)
276 all.MSPEs[i, "Min-cv tree"] = MSPE.fit.tree.min
277
278
279 # "1-se tree"
280 err.min = info.tree[ind.min, "xerror"]
281 se.min = info.tree[ind.min, "xstd"]
282 threshold = err.min + se.min
283 ind.lse = min(which(info.tree[1:ind.min, "xerror"] < threshold))
284
285 ### Get the corresponding CP value, averaging if necessary
286 CP.lse.raw = info.tree[ind.lse, "xerror"]
287 if(ind.lse == 1){
288     ### If best CP is in row 1, store this value
289     CP.lse = CP.lse.raw
290 } else{
291     ### If best CP is not in row 1, average this with the value from the
292     ### row above it.

```



```

293
294     ### Value from row above
295     CP.above = info.tree[ind.lse-1, "CP"]
296
297     ### (Geometric) average
298     CP.lse = sqrt(CP.lse.raw * CP.above)
299 }
300
301 fit.tree.lse = prune(fit.tree, cp = CP.lse)
302 fit.tree.lse.pred = predict(fit.tree.lse, data.valid)
303 MSPE.fit.tree.lse = get.MSPE(Y.valid, fit.tree.lse.pred)
304 all.MSPEs[i, "1-se tree"] = MSPE.fit.tree.lse
305 }
306
307 all.MSPEs
308
309 ### Make a boxplot of MSPEs. I would like to include the number of folds
310 ### in the title. This can be done by using the paste0() function,
311 ### which concatenates strings (i.e. attaches them end-to-end), and
312 ### can be provided numeric variables.
313 boxplot(all.MSPEs, main = paste0("CV MSPEs over ", K, " folds"))
314
315
316
317 ### Calculate RMSPEs
318 all.RMSPEs = apply(all.MSPEs, 1, function(w){
319     best = min(w)
320     return(w / best)
321 })
322 all.RMSPEs = t(all.RMSPEs)
323
324 ### Make a boxplot of RMSPEs
325 boxplot(all.RMSPEs, main = paste0("CV RMSPEs over ", K, " folds"))
326
327 > all.MSPEs

```

	LS	Hybrid	Ridge	LASSO-Min	LASSO-lse	GAM	Full-tree	Min-cv tree	1-se tree
[1,]	1037.2843	935.5934	1026.2063	1031.8530	1081.5330	135.46137	521.42750	531.63564	1103.9618
[2,]	419.5374	520.7976	462.0808	463.0394	426.9839	1166.67426	129.56077	122.76918	598.1703
[3,]	561.5317	606.2170	569.6294	559.5286	427.4034	674.21375	209.19630	209.19630	284.9489
[4,]	117.0370	140.0399	119.3903	119.7927	194.3578	140.95548	98.93844	98.93844	199.0212
[5,]	200.1188	195.1213	201.8795	230.7676	285.8414	42.07896	115.41266	112.12023	458.6532
[6,]	208.1940	180.5834	208.5333	209.9717	220.0076	297.97684	386.37579	386.37579	752.4572
[7,]	507.6231	555.5441	515.8408	654.4310	1062.9010	231.24568	158.20856	163.34328	305.6825
[8,]	151.2711	150.9220	149.6126	150.8763	319.8533	199.88708	234.91171	254.13553	336.9521
[9,]	681.4810	731.9910	687.8407	749.9537	928.2086	854.82298	418.22831	421.05592	434.6720
[10,]	358.1365	256.2422	365.3990	398.0162	281.7134	262.95811	218.79173	218.79173	242.9717

```

327 mean(all.MSPEs[,7])
328 mean(all.MSPEs[,8])
329 mean(all.MSPEs[,9])

```

(a) Compute the mean MSPE for each tree model and comment on the comparison.


```

> mean(all.MSPEs[,7])
[1] 249.1052
> mean(all.MSPEs[,8])
[1] 251.8362
> mean(all.MSPEs[,9])
[1] 471.7491

```

Full-tree and Min-cv tree has similar small MSPEs and 1-se tree has quite larger MSPE.

(b) ADD the three trees to the relative MSPE boxplots made previously.
Comment on how well they perform compared to other methods.



So far, it looks like Full-tree and Min-cv tree and GAM show the best result. However, 1-se tree doesn't do well on it.